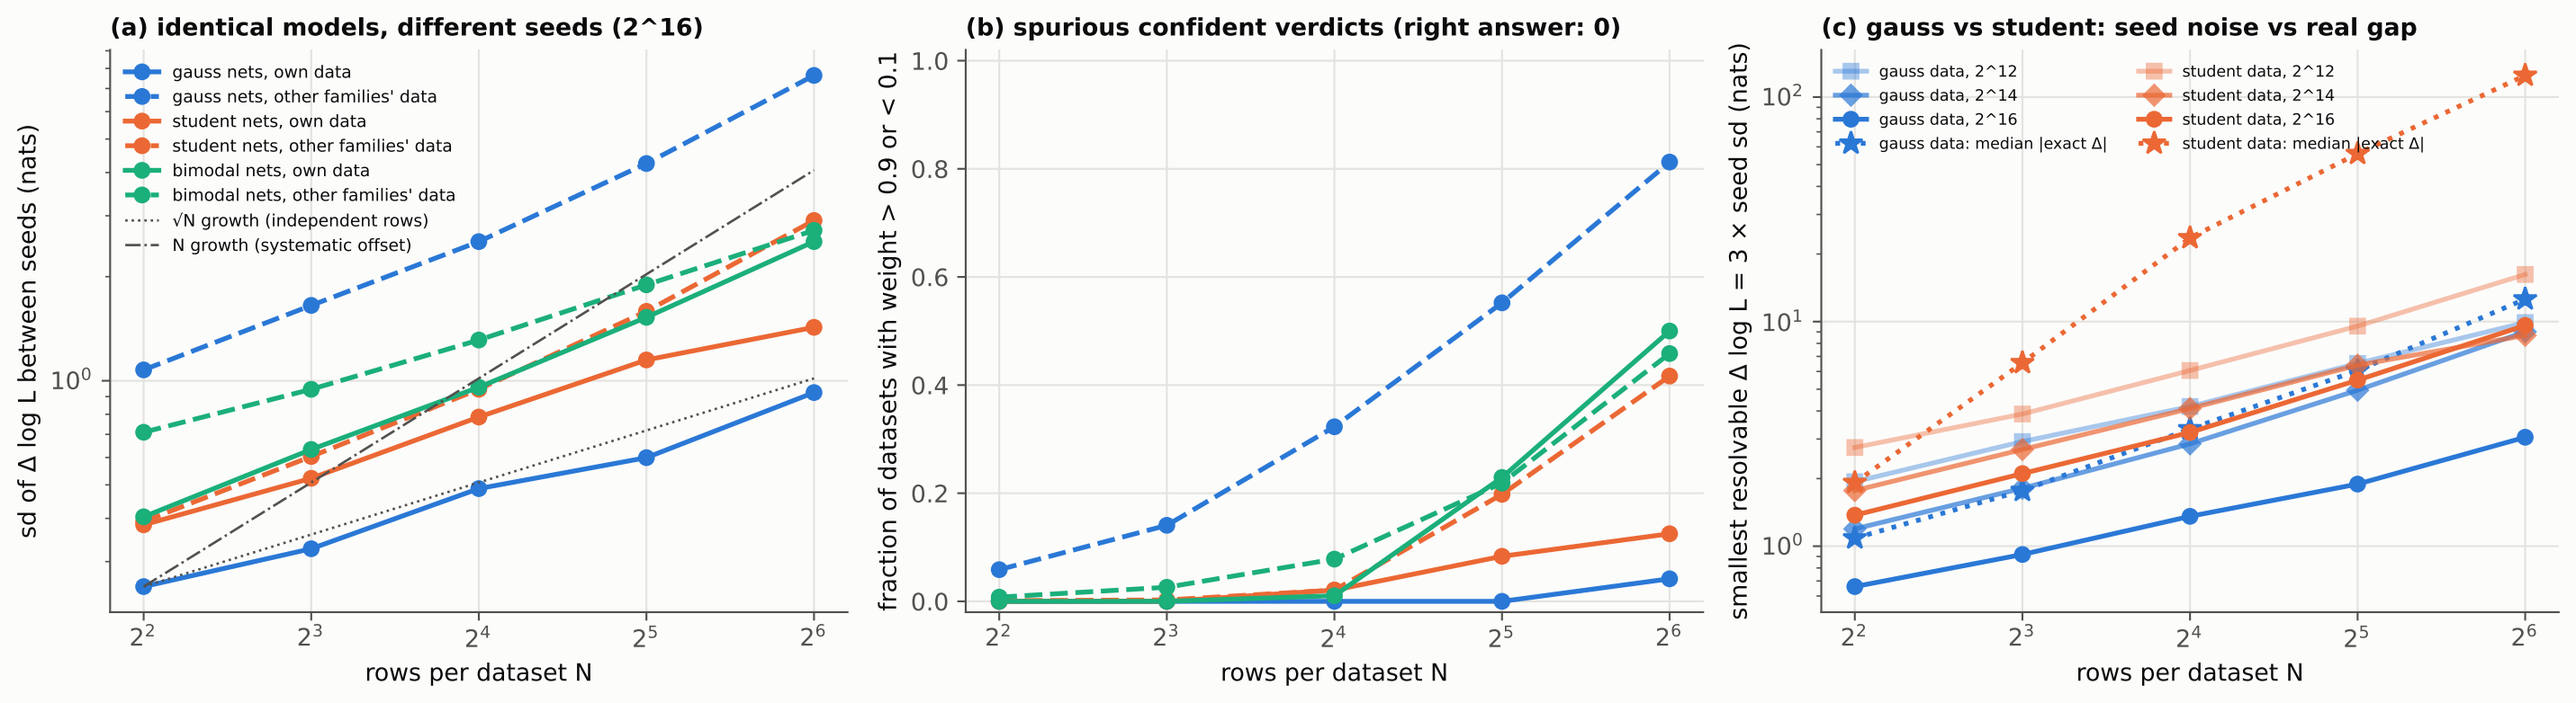
(a) identical models, different seeds (2^16)
gauss nets, own data
gauss nets, other families' data
student nets, own data
student nets, other families' data
bimodal nets, own data
bimodal nets, other families' data
√N growth (independent rows)
N growth (systematic offset)
sd of Δ log L between seeds (nats)
rows per dataset N
(b) spurious confident verdicts (right answer: 0)
fraction of datasets with weight > 0.9 or < 0.1
rows per dataset N
(c) gauss vs student: seed noise vs real gap
gauss data, 2^12
gauss data, 2^14
gauss data, 2^16
gauss data: median |exact Δ|
student data, 2^12
student data, 2^14
student data, 2^16
student data: median |exact Δ|
smallest resolvable Δ log L = 3 × seed sd (nats)
rows per dataset N