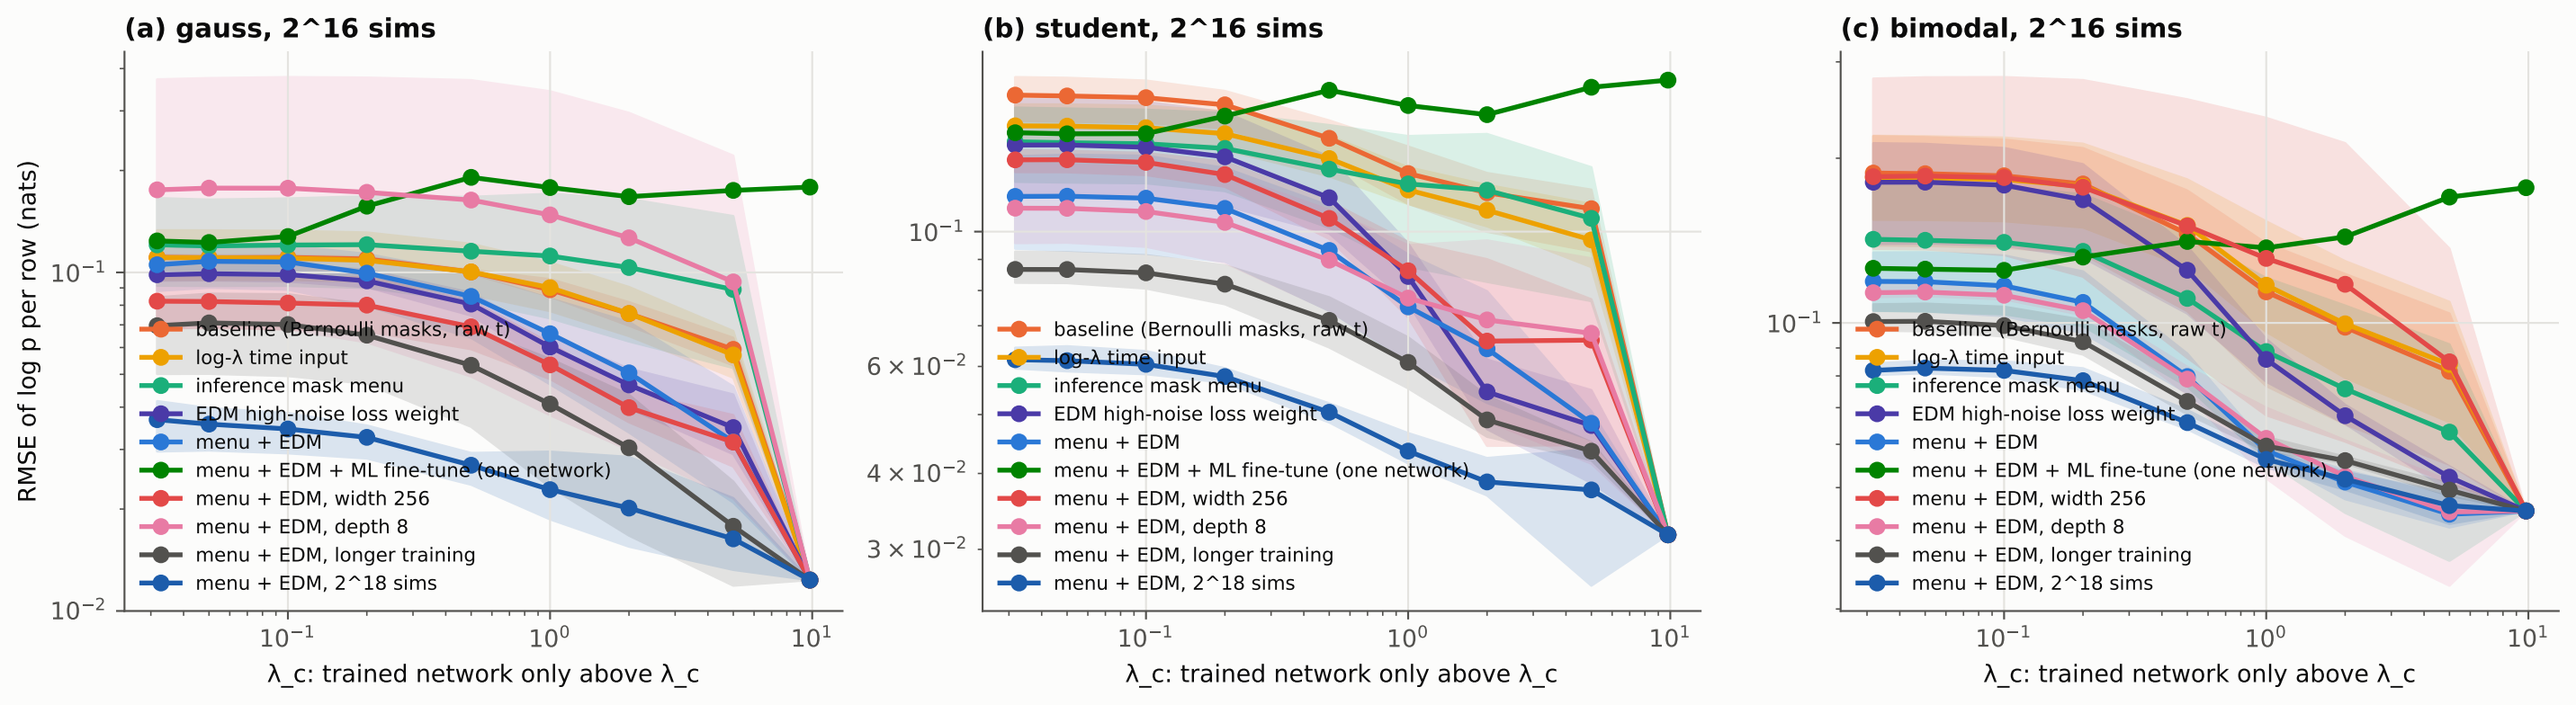
(a) gauss, 2^16 sims
(b) student, 2^16 sims
(c) bimodal, 2^16 sims
RMSE of log p per row (nats)
λ_c: trained network only above λ_c
baseline (Bernoulli masks, raw t)
log-λ time input
inference mask menu
EDM high-noise loss weight
menu + EDM
menu + EDM + ML fine-tune (one network)
menu + EDM, width 256
menu + EDM, depth 8
menu + EDM, longer training
menu + EDM, 2^18 sims